

# REPAIR: ROBUST EDITING VIA PROGRESSIVE ADAPTIVE INTERVENTION AND REINTEGRATION

Anonymous authors

Paper under double-blind review

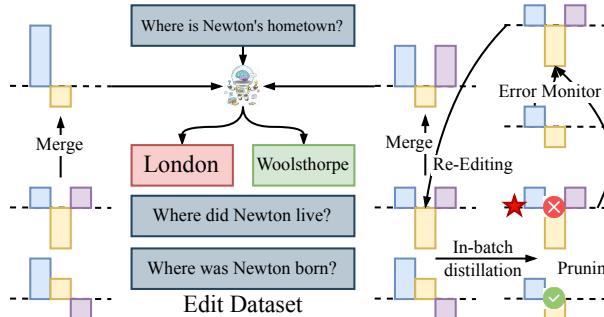
## ABSTRACT

Post-training for large language models (LLMs) is constrained by the high cost of acquiring new knowledge or correcting errors and by the unintended side effects that frequently arise from retraining. To address these issues, we introduce REPAIR (Robust Editing via Progressive Adaptive Intervention and Reintegration), a lifelong editing framework designed to support precise and low-cost model updates while preserving non-target knowledge. REPAIR mitigates the instability and conflicts of large-scale sequential edits through a closed-loop feedback mechanism coupled with dynamic memory management. Furthermore, by incorporating frequent knowledge fusion and enforcing strong locality guards, REPAIR effectively addresses the shortcomings of traditional distribution-agnostic approaches that often overlook unintended ripple effects. Our experiments demonstrate that REPAIR boosts editing accuracy by 10%-30% across multiple model families and significantly reduces knowledge forgetting. This work introduces a robust framework for developing reliable, scalable, and continually evolving LLMs.

## 1 INTRODUCTION

Large language models (LLMs) have demonstrated remarkable capabilities across diverse tasks. However, their inherent rigidity prevents them from autonomously updating knowledge after pre-training, rendering them unable to correct errors (*e.g.*, hallucinations or outdated facts) or integrate new information. Consequently, lifelong model editing has emerged as a critical research paradigm. It aims to enable continuous, efficient, and low-cost local updates that ensure models remain accurate and relevant over time Wang et al. (2024b). In contrast to full re-training or broad fine-tuning, editing focuses on *precisely fine-grained* modifications that preserve unrelated competencies while delivering immediate corrections at deployment time.

Despite steady progress, important gaps remain as shown in Figure 1. **(1) Large-scale sequential editing & coarse knowledge fusion.** As edits accumulate, models can exhibit routing instability, conflicts among edits, and even collapse. Thus, stabilizing sequential updates without broad side effects remains challenging Gupta et al. (2024); Cohen et al. (2024). Semi-parametric designs (*e.g.*, SERAC Mitchell et al. (2022b)) and discrete key-value adaptors (*e.g.*, GRACE Hartvigesen et al. (2023)) alleviate some failure modes and support long edit streams, but still face scope and auditing trade-offs Mitchell et al. (2022b); Hartvigesen et al. (2023). The strategy for knowledge fusion remains underexplored, despite being the stage most prone to information loss Wang et al. (2024a). **(2) Few-shot editing.** Under data-scarce conditions, editors often struggle to form robust, generalizable



**Figure 1: Problems and our solutions.** REPAIR achieves closed-loop feedback, fine-grained knowledge integration, weighted knowledge merging, and robust editing performance.

054 changes beyond the exact prompt, motivating gradient-transformation editors trained for locality  
 055 (*e.g.*, MEND Mitchell et al. (2022a)) and broader taxonomies of edit generalization Mitchell et al.  
 056 (2022a); Wang et al. (2024b). **(3) Open-loop and distribution-agnostic learning.** Many pipelines  
 057 operate without reflective feedback, optimize on indiscriminate batches, and under-stress-test ripple  
 058 effects on related knowledge and reasoning, calling for tighter evaluation and integration mechanisms Cohen et al. (2024); Wang et al. (2024b). Overall, these issues reveal a fundamental trade-off  
 059 among reliability, specificity, and scalability that any practical editing system must reconcile.  
 060

061 To address these challenges, we propose the framework named REPAIR (**R**obust **E**diting via  
 062 **P**rogressive **A**daptive **I**ntervention and **R**eintegration), with targeted strategies: **(1) Closed-loop**  
 063 **feedback with dynamic memory management** that monitors edit performance and selectively re-  
 064 initializes underperforming modules to stabilize routing and consolidation at scale. Concretely, our  
 065 controller triggers health checks after each edit window and performs scoped resets or compaction  
 066 when drift is detected. **(2) Distribution-aware optimization** that reorganizes samples by similar-  
 067 ity and applies inner-batch distillation to enhance consistency and robustness in few-shot settings,  
 068 encouraging edits to generalize across paraphrases and nearby contexts rather than overfitting to sin-  
 069 gle prompts. **(3) Frequent knowledge fusion** that increases fusion cadence to prevent information  
 070 loss and ensure timely consolidation of new and existing knowledge, with guardrails that validate  
 071 locality before integration to avoid unintended side effects.

072 We compare REPAIR with several foundational model editing methods across three dimensions:  
 073 *Memory*, *Attributes*, and *Behaviors* (Table 1). Its core innovation lies in integrating a dual memory  
 074 system with parametric editing, complemented by error feedback, inner-batch knowledge distilla-  
 075 tion, and loss-aware subspaces merging. This design achieves high success rates and broad editing  
 076 coverage while minimizing side effects. In contrast, previous methods struggle with knowledge  
 077 overlap and loss, particularly in sequential editing, where large differences between adjacent sam-  
 078 ples hinder effective correction. Table 2 showcases cases where REPAIR outperforms baselines,  
 079 offering a better balance of Reliability, Generalization, and Locality.

080 **Table 1: Comparison of current model editing methods.** “✓” refers to “yes” and “well-  
 081 supported”, “✗” refers to “no” or “badly-supported”, and “○” refers to “less-supported”. The three  
 082 metrics of Reliability, Generalization, and Locality denote the performance on lifelong editing.

Methods	Memory			Attributes			Behaviors		
	Long-term Memory	Working Memory	Parametric	Lifelong	Reliability	Generalization	Locality	Error Feedback	Knowledge Distillation
FT-EWC Kirkpatrick et al. (2017)	✓	✗	✓	✓	✓	✓	✗	✗	✗
ROME Meng et al. (2022b)	✓	✗	✓	✗	✗	✗	✗	✗	✗
MEMIT Meng et al. (2023)	✓	✗	✓	✗	✗	✗	✗	✗	✗
MEND Mitchell et al. (2022a)	✓	✗	✓	✗	✗	✗	✗	✗	✗
DEFER Mitchell et al. (2022b)	✗	✓	✓	✓	○	✗	✗	✗	✗
GRACE Hartvigsen et al. (2023)	✗	✓	✗	✓	✓	✗	✓	✗	✗
WISE Wang et al. (2024a)	✓	✓	✓	✓	✓	✓	✓	✗	✗
REPAIR	✓	✓	✓	✓	✓	✓	✓	✓	✓

093 **Table 2: Failure cases study.** Previous baselines(Wang et al. (2024a)Hartvigsen et al. (2023))often  
 094 encounter issues of repeating answers from previous questions and difficulty in correcting adjacent  
 095 knowledge during editing.

MethodPrompt	Edit Target	Post-Edit Output	Metrics
a) The genus <i>Platypatrobis</i> is part of the family?	Arctiinae	Arctiue ✗	Reliability ✗
b) <i>The genus Platypatrobis is a part of what family</i>	-	Yemen ✗	Generalization ✗
c) <i>The genus Platypatrobis is part of the family?</i>	-	Arctiinae ✓	
c) When was the IAAF Combined Events Challenge launched?	2006	Armand ✗	Reliability ✗
d) When does season 5 of ruby come out?	October 14, 2017	2006 ✗	Locality ✗
e) <i>when does season 5 of ruby come out?</i>	-	2017 ✓	

104 In summary, the main contributions are as follows.  
 105

- 106 • We identify three critical challenges in model editing: (1) instability under large-scale se-  
 107 quential edits, (2) limited generalization in few-shot scenarios, and (3) inefficiency in open-  
 108 loop, distribution-agnostic pipelines.

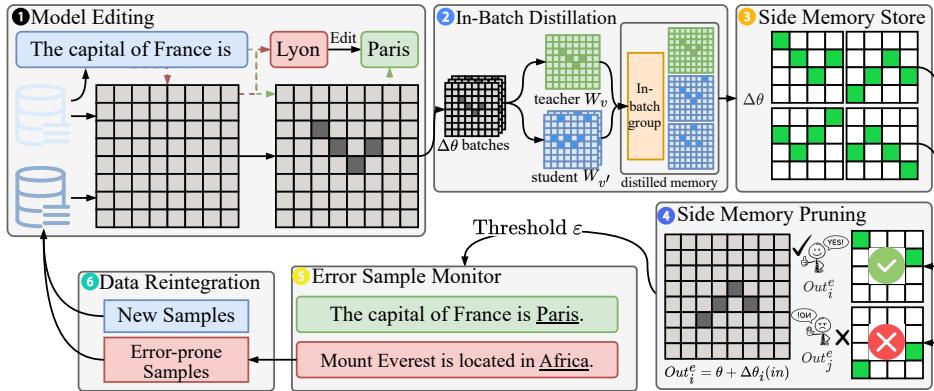


Figure 2: **The overall structure of REPAIR.** An edit, such as changing the capital of France from "Lyon" to "Paris," is stored as a parameter update,  $\Delta\theta$ , in the Side Memory. An Error Sample Monitor evaluates the performance of each edit ( $Out_i^e$ ). If the error rate,  $Err_{thresh}$ , for an edit on a new sample exceeds a threshold  $\epsilon$ , the Side Memory Pruning module removes the erroneous update. The system then reintegrates new and error-prone samples for continuous learning.

- We propose REPAIR, a novel framework to address these challenges by integrating a dual-memory system with parametric editing. It introduces closed-loop error feedback, distribution-aware optimization, and loss-aware subspaces merging to ensure robust and precise updates.
- We validate the performance of REPAIR across diverse models (including LLaMA-3, Qwen-2.5, DeepSeek-R1-1.5B, and GPT-2-XL), demonstrating a 15%–20% improvement in overall editing performance over state-of-the-art methods and showing consistent, robust generalization.

## 2 METHODOLOGY

We propose a novel closed-loop lifelong model editing framework, denoted as **REPAIR**, which addresses the limitations of open-loop editing in distributed side-memory methods. Our framework, as shown in Figure 2, integrates (1) closed-loop error feedback with dynamic memory management; (2) distribution-aware batch reassembly with inner-batch knowledge distillation; (3) loss-aware weighted knowledge merging.

### 2.1 PROBLEM SETUP

**Definition 2.1** (Lifelong Model Editing). *Given a pre-trained model  $f_{\theta_0}(y|x)$ , a sequential edit stream  $\{\mathcal{E}_t\}_{t=1}^T$  where  $\mathcal{E}_t = \{(x_i^{(t)}, y_i^{(t)})\}_{i=1}^N$ , and auxiliary distributions  $\mathcal{G}(x)$  (paraphrased inputs) and  $\mathcal{U}$  (unrelated contexts), the objective is to obtain updated parameters  $\theta_T$  that optimize the multi-objective trade-off:*

$$\theta_t = \arg \min_{\theta} \alpha \underbrace{\frac{1}{N} \sum_{i=1}^N \ell \left( f_{\theta}(\cdot|x_i^{(t)}), y_i^{(t)} \right)}_{\text{reliability}} + \beta \underbrace{\frac{1}{N} \sum_{i=1}^N \mathbb{E}_{x' \sim \mathcal{G}(x_i^{(t)})} \left[ \ell \left( f_{\theta}(\cdot|x'), y_i^{(t)} \right) \right]}_{\text{generalization}} \\ + \gamma \underbrace{\mathbb{E}_{x \sim \mathcal{U}} \left[ \text{KL} \left( f_{\theta_{t-1}}(\cdot|x) \| f_{\theta}(\cdot|x) \right) \right]}_{\text{locality}} + \underbrace{R(\theta, \theta_{t-1})}_{\text{stability}} \quad (1)$$

where  $(\alpha, \beta, \gamma)$  are hyperparameters controlling the reliability-generalization-locality-stability trade-off, and  $R$  denotes a regularization term enforcing parameter smoothness across sequential edits.

162    2.2 DUAL MEMORY MECHANISM AND ROUTING  
 163

164    As shown in Figure 2, block 1: For dual memory-based editing methods, the dual memory mecha-  
 165    nism is typically deployed in the deep layers of the network. Specifically, for the value matrix  $\mathbf{W}_v$   
 166    of the target FFN layer, here create a copy as the side memory pool  $M_s$ , i.e.:  $M_s^{(0)} = W_v$ . If the side  
 167    memory pool is activated, the output is computed as:  $o_s = \phi(f^T W_k) \cdot M_s$ , where  $\phi$  denotes the  
 168    non-linear activation function, and  $o_s$  represents the FFN output based on the side memory(Wang  
 169    et al. (2024a)).

170    During the inference phase, for moemory pool  $i$ , the activation score is defined as  
 171

$$\Delta_{\text{act}}^{(i)}(x) = \|\mathcal{A}(x) \cdot (W'_{v,i} - W_v)\|_2. \quad (2)$$

172    where  $\mathcal{A}(\cdot) = a$  is the activation of the side memory’s corresponding FFN layer. Routing selects the  
 173    pool with the activation score. If  $\max_i \Delta_{\text{act}}^{(i)}(x) \leq \varepsilon$ , the main memory  $W_v$  is used. Otherwise, the  
 174    side memory pool  $M_s$  is selected. To enforce discriminative routing, we use a margin-based loss.  
 175    The objective of the routing mechanism is to establish a clear decision boundary:  
 176

$$\min_{\mathbf{x}_e \sim \mathcal{E}} \mathcal{R}(\mathbf{x}_e) \sim \min_{\mathbf{x}' \sim \mathcal{U}} \mathcal{R}(\mathbf{x}') > \tau > \max_{\mathbf{x}_i \sim \mathcal{G}} \mathcal{R}(\mathbf{x}_i) \quad (3)$$

177    where  $\tau$  is a preset threshold, and  $\mathcal{E}$  and  $\mathcal{G}_i$  represent the edit and edit-irrelevant datasets, respec-  
 178    tively. This selective activation mechanism ensures that edited knowledge is only retrieved in rele-  
 179    vant contexts, thereby minimizing interference with the original model’s performance.

180    2.3 DISTRIBUTION-AWARE INNER-BATCH KNOWLEDGE DISTILLATION  
 181

182    As shown in Figure 2 block 2: A sample batch  $\mathcal{E} = \{x_1, x_2, \dots, x_n\}$ , and denote the corresponding  
 183    feature representations by  $o_i = \text{Norm}(f_\theta(x_i))$ ,  $i = 1, \dots, n$ . To improve the consistency and  
 184    stability of model updates during sequential edits, we organized samples into homogeneous batches  
 185    and performed intrabatch knowledge distillation. Samples with high mutual similarity are grouped  
 186    into a batch  $B = \{x^{(0)}, x^{(1)}, \dots, x^{(b-1)}\}$ . Within each batch, the first sample  $x^{(0)}$  acts as a *teacher*,  
 187    while the remaining samples are *students*. We define the inner-batch knowledge distillation loss as  
 188

$$\mathcal{L}_{\text{KD}} = \lambda \cdot \mathcal{L}_{\text{cosine}} + \theta \cdot \mathcal{L}_{\text{variance}} \quad (4)$$

189    where  $\mathcal{L}_{\text{cosine}} = 1 - \frac{o_i \cdot o_0}{\|o_i\| \|o_0\|}$  and  $\mathcal{L}_{\text{variance}} = \frac{1}{N} \sum_{i=1}^N \|o_i - o_{\text{mean}}\|^2$ . Minimizing  $\mathcal{L}_{\text{kd}}$  encourages  
 190    all samples in the batch to share similar knowledge, which in turn reduces potential conflicts when  
 191    updating the same network parameters  $\theta$ . The regularization term is used to maintain diversity  
 192    among features, preventing excessive uniformity.

193    If certain samples cannot be well-aligned with the batch (i.e., their  $\mathcal{L}_{\text{kd}}$  remains high after optimiza-  
 194    tion), this indicates that they do not belong to the same distribution cluster and are unlikely to be  
 195    effectively edited together. Such samples are removed from the batch and reclustered with other  
 196    samples to form new homogeneous groups. Formally, the final batch reassembly can be expressed  
 197    as

$$\mathcal{B}^* = \text{Recluster}(\{x \in B \mid \mathcal{L}_{\text{kd}}(x, B) < \epsilon\}), \quad (5)$$

198    where  $\epsilon$  is a threshold controlling inner-batch consistency. This procedure ensures that sequential  
 199    parameter edits are performed on groups of samples with aligned knowledge, improving both sta-  
 200    bility and effectiveness of the model update. The convergence proof is provided in the Appendix 4  
 201    and Appendix 2.

202    2.4 CLOSED-LOOP ERROR FEEDBACK AND MEMORY PRUNING  
 203

204    As shown in Figure 2 block 4: After each editing cycle, we evaluate the performance in a feedback  
 205    pool  $\mathcal{E}$  of error response samples by comparing to the correctness threshold  $\tau_{\text{correct}}$ . For each shard  
 206     $i$ , we define the error set  $\mathcal{E}_i = \{x \in \mathcal{E} \mid i^*(x) = i\}$  and compute the error rate  $r_i^{\text{pool}}$  for each  
 207    side memory pool, defined as the proportion of failed edits within the corresponding sample set:  

$$r_i^{\text{pool}} = \frac{|\{x \in \mathcal{E}_i \mid a(x) \leq \tau_{\text{correct}}\}|}{|\mathcal{E}_i|}$$

208    When the pruning conditions are met ( $r_i > \tau_{\text{prune}}$  or  $|\mathcal{E}| > \tau_E$ ), we execute the following procedure:

- 216     1. **Memory pool screening & pruning:** Identify the side memory pool with the highest error  
 217       rate  $j = \arg \max_i r_i^{\text{pool}}$ . Remove the identified memory pool from the system.  
 218  
 219     2. **Sample Reintegration & retraining:** Recombine the remaining error samples to form a  
 220       new training set  $\mathcal{E}_{\text{retrain}}$ . Retrain the new side memory pools using  $\mathcal{E}_{\text{retrain}}$ .

221     This closed-loop feedback mechanism enables the system to dynamically identify and eliminate  
 222       underperforming memory units while optimizing the overall editing performance through sample  
 223       reorganization and iterative retraining. The time-convergence proof is provided in the Appendix 2.  
 224

## 225     2.5 MERGING WITH WEIGHTED TIES

227     As shown in Figure 2 block 3: After multiple updates, shards  $\{W'_{v,i}\}$  produce deltas  $\tau_i = W'_{v,i} - W_v$ .  
 228     We merge them with the weighted TIES Yadav et al. (2023) operator based on :  $W'_v \leftarrow W_v +$   
 229        $\omega_i$  TIES  $(\{\tau_i\}_{i=1}^k; W_v)$ .

230     The total loss integrates all components:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{edit}} + \lambda_a \mathcal{L}_a + \lambda_{\text{KD}} \mathcal{L}_{\text{KD}}. \quad (6)$$

233      $\mathcal{L}_{\text{edit}}$  is the autoregressive cross-entropy.  $\mathcal{L}_{\text{edit}}(W'_v) = -\log P_{W'_v}(y | x)$ . To enforce discriminative  
 234       routing, we use a margin-based loss:

$$\begin{aligned} \mathcal{L}_a = \min \Bigg\{ & \max(0, \Delta_{\text{act}}(x_i) - \gamma_1) \\ & + \max(0, \gamma_2 - \Delta_{\text{act}}(x_e)) + \max(0, \gamma - (\Delta_{\text{act}}(x_e) - \Delta_{\text{act}}(x_i))) \Bigg\} \end{aligned} \quad (7)$$

240     For shard  $i$ , consider  $\parallel$  subspaces  $\{\theta_1, \dots, \theta_k\}$ , each trained on a subset of samples  $\mathcal{E}_i$ . Let the  
 241       average training loss of subspaces  $\theta_i$  be:  $\mathcal{L}_i = \frac{1}{|\mathcal{E}_i|} \sum_{(x,y) \in \mathcal{E}_i} \ell(f(x; \theta_i), y)$ , where  $\ell(\cdot)$  is the task  
 242       loss. We define the merging weight of each subspaces as  $w_i = \frac{\exp(-\alpha \mathcal{L}_i)}{\sum_{j=1}^M \exp(-\alpha \mathcal{L}_j)}$ , with  $\alpha > 0$   
 243       controlling sensitivity to the loss. The global network parameters are then obtained via weighted  
 244       averaging:  $\theta = \sum_{i=1}^M w_i \theta_i$ . This loss-aware merging favors subspaces that achieve lower training  
 245       loss on their corresponding samples, promoting reliable knowledge integration.  
 246

## 247     3 EXPERIMENTS

250     In the experimental section, we design six evaluations to answer the following questions: **Q1**, do  
 251       the three key innovations (closed-loop feedback, discriminative pruning, and distribution reinte-  
 252       gration) improve edit accuracy, generalization, and locality? **Q2**, does the method generalize well  
 253       to knowledge-intensive tasks such as question answering and hallucination mitigation? **Q3**, is the  
 254       method effective across different parameter scales and diverse architectures, including recent open-  
 255       source models? **Q4**, under distribution shift (e.g., on the Wikibig Edit dataset), does the method  
 256       remain robust and outperform existing methods? **Q5**, can the method maintain long-term stabili-  
 257       ty and reliability in large-scale sequential editing scenarios? **Q6**, what are the contributions and  
 258       sensitivities of each component and hyperparameter to overall performance?

### 259     3.1 EXPERIMENTAL SETUP

261     **Datasets and Models.** Autoregressive LLMs are ideal for evaluating model editing due to their  
 262       unidirectional causal structure, which allows predictable and traceable edits. This ensures clear  
 263       interpretability of edit generalization and locality. We evaluate widely used models (LLaMA-3-  
 264       8B, GPT2-XL) and recent models (Qwen2.5-7B, DeepSeek-R1-1.5B), using datasets such as ZsRE  
 265       for closed-book QA, Wikibig Edit for editing performance, and a hallucination dataset to assess  
 266       generalization. For more details, refer to the Appendix 5.

### 267     Baselines.

- 269       • **Direct Parameter Editors:** Directly modify model weights (e.g., **ROME** Gupta et al.  
                  (2024), **MEMIT** Meng et al. (2023), **MEMIT-mass** Meng et al. (2023)).

- 270  
 271  
 272  
 273  
 274  
 275
- **Hypernetwork-Based Editors:** Use an auxiliary network to generate parameter updates at inference (e.g., **MEND** Mitchell et al. (2022a)).
  - **External Memory-Based Editors:** Leave the model unchanged and store edits in external memory, retrieved via a routing mechanism (e.g., **SERAC** Mitchell et al. (2022b), **GRACE** Hartvigsen et al. (2023), **WISE** Wang et al. (2024b)).

276 **Implementation Details.** experiments were conducted simultaneously using two GPUs: an A100  
 277 PCIe 80GB and an A100 SXM4 40GB. The code was implemented based on PyTorch 2.1, with  
 278 modifications built upon the original EasyEditor framework. The specific hyperparameter settings  
 279 are detailed in Appendix C.

280 **Evaluation Metrics** Each edited corpus instance comprises three components: the descriptor  $k_e$   
 281 used to perform the edit, an irrelevant prompt-answer pair  $k'_e$  to verify locality and a rephrase prompt  
 282  $k_{loc}$  to evaluate generalization performance across different expressions. To comprehensively eval-  
 283 uate the optimization capability of the proposed method in addressing the continual learning trilemma,  
 284 we employ four metrics—edit accuracy:  $\text{Rel} = \frac{1}{N} \sum_{n=1}^N \mathbb{I}(f_{\omega_N}(\mathbf{x}_e^n) = \mathbf{y}_e^n)$ , rephrase accuracy :  
 285  $\text{Gen} = \frac{1}{N} \sum_{n=1}^N \mathbb{I}(f_{\omega_N}(\mathbf{x}'_e^n) = \mathbf{y}'_e^n)$ , locality :  $\text{Loc} = \frac{1}{N} \sum_{n=1}^N \mathbb{I}(f_{\omega_N}(\mathbf{x}_{loc}) = f_{\omega_0}(\mathbf{x}_{loc}))$ . We use  
 286 the geometric mean of Rel., Gen., and Loc. to evaluate the overall editing performance, which bal-  
 287 ances metric sensitivity and interpretability, exhibits sensitivity to weak performance areas, and is  
 288 suitable for scenarios where all three metrics are equally important.  $\text{OP} = \sqrt[3]{\text{Rel.} \times \text{Gen.} \times \text{Loc.}}$  to  
 289 assess the holistic editing effectiveness. Here,  $\mathbb{I}(\cdot)$  is the indicator function used to count the number  
 290 of successful predictions.

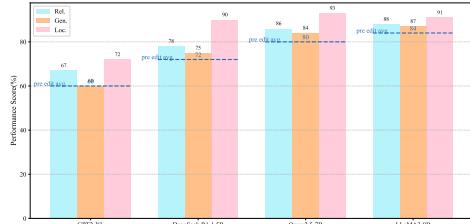
291 For the hallucination dataset specifically, we utilize perplexity(PPL) as the metric to assess editing  
 292 performance. PPL can be interpreted as the "average branching factor in predicting the next token,"  
 293 where a lower value indicates more accurate model predictions and suggests a reduced likelihood of  
 294 the edited model generating hallucinations.  $\text{PPL} = \exp\left(-\frac{1}{N} \sum_{i=1}^N \log P(y_i|\text{context}_i)\right)$

### 297 3.2 MAIN RESULTS

298 Table 3 effevtively addressed **Q1**, **Q4** and **Q5**. It has been rigorously evaluated across diverse  
 299 models and scales ( $N = 1, 30, 120, 1000$ )  
 300 of QA editing tasks, demonstrating state-of-  
 301 the-art performance. Fine-tuning-based meth-  
 302 ods achieve good accuracy and generalization  
 303 at small scales but suffer from catastrophic  
 304 forgetting and knowledge conflicts in large-  
 305 scale edits, leading to performance degradation.  
 306 GRACE excels in accuracy but has limited  
 307 generalization, while WISE maintains strong local-  
 308 ity but sacrifices critical knowledge, reducing  
 309 editing accuracy. ROME-style methods are stable but overfit and struggle with generalization.

310 To address **Q2**, Table 4 shows REPAIR’s effectiveness in reducing hallucinations on the SelfCheck-  
 311 GPT dataset for LLaMA-3-8B across different editing scales. REPAIR balances reduced hallucina-  
 312 tions with preserved locality, making it highly effective for large-scale model editing.

313 To address **Q3** and **Q4**, Table 3 and Figure 3 show that REPAIR’s closed-loop error feedback,  
 314 together with distribution-aware clustering and redistribution, yields consistently superior perfor-  
 315 mance across edit scales and exceptional stability for large-scale edits. Smaller models concen-  
 316 trate knowledge in narrower parameter subsets, enabling reliable local corrections but weaken-  
 317 ing long-term stability and generalization (i.e., maintaining accuracy while preserving unrelated  
 318 knowledge). Accordingly, DeepSeek-R1-1.5B attains higher immediate correction rates at small  
 319 edit\_Num, yet degrades quickly as N grows. For locality, LLaMA-3-8B and Qwen2.5-7B are  
 320 marginally stronger due to parameter redundancy; DeepSeek-R1-1.5B remains competitive only at  
 321 low N, then collapses under extreme multi-point editing. In contrast, larger models distribute knowl-  
 322 edge more broadly, and though harder to modify—successful edits generalize better across contexts.  
 323 At a medium scale ( $N=120$ ), MEMIT-M and WISE show higher Rel., likely because REPAIR’s



324 Figure 3: Average Editing Performance of Wik-  
 325 iBigEdit Across Different Models

326

327

328

329

330

331

332

333

334

335

336

337

338

339

340

341

342

343

344

345

346

347

348

349

350

351

352

353

354

355

356

357

358

359

360

361

362

363

364

365

366

367

368

369

370

371

372

373

374

375

376

377

378

379

380

381

382

383

384

385

386

387

388

389

390

391

392

393

394

395

396

397

398

399

400

401

402

403

404

405

406

407

408

409

410

411

412

413

414

415

416

417

418

419

420

421

422

423

424

425

426

427

428

429

430

431

432

433

434

435

436

437

438

439

440

441

442

443

444

445

446

447

448

449

450

451

452

453

454

455

456

457

458

459

460

461

462

463

464

465

466

467

468

469

470

471

472

473

474

475

476

477

478

479

480

481

482

483

484

485

486

487

488

489

490

491

492

493

494

495

496

497

498

499

500

501

502

503

504

505

506

507

508

509

510

511

512

513

514

515

516

517

518

519

520

521

522

523

524

525

526

527

528

529

530

531

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

550

551

552

Table 3: Comparative results for QA on multi-scale editing (ZsRE and WikiBigEdit)  $N$ : Num Edits.

Method	$N = 1$			$N = 30$			$N = 120$			$N = 1000$		
	Rel.	Gen.	Loc.	OP.	Rel.	Gen.	Loc.	OP.	Rel.	Gen.	Loc.	OP.
LLaMA-3-8B (ZsRE)												
FT-L	0.57	0.52	0.96	0.66	0.35	0.35	0.52	0.39	0.29	0.26	0.21	0.25
FT-EWC	0.96	<b>0.93</b>	0.02	0.26	0.78	0.76	0.02	0.23	0.76	<b>0.76</b>	0.08	0.36
MEND	0.95	0.93	0.96	0.95	0.24	0.25	0.18	0.22	0.08	0.07	0.00	0.00
ROME	0.85	0.80	0.99	0.88	0.61	0.60	0.68	0.63	0.22	0.22	0.04	0.12
MEMIT-M	0.84	0.81	0.99	0.88	0.73	0.72	0.95	0.79	0.70	0.65	0.82	0.72
DEFER	0.68	0.58	0.56	0.61	0.65	0.47	0.36	0.49	0.20	0.12	0.27	0.20
GRACE	<b>0.97</b>	0.36	<b>1.00</b>	0.71	<b>0.96</b>	0.17	<b>1.00</b>	0.55	<b>0.94</b>	0.14	<b>1.00</b>	0.51
WISE	0.94	0.92	<b>1.00</b>	<b>0.95</b>	0.62	0.60	0.86	0.68	0.57	0.58	0.87	0.66
<b>REPAIR</b>	0.94	0.92	<b>1.00</b>	<b>0.95</b>	0.93	<b>0.90</b>	0.87	<b>0.89</b> ↑	0.76	0.74	<b>1.00</b>	<b>0.83</b> ↑
Qwen2.5-7B (ZsRE)												
FT-L	0.68	0.63	0.93	0.74	0.28	0.23	0.44	0.30	0.13	0.11	0.10	0.11
FT-EWC	0.97	0.92	0.05	0.35	0.82	0.80	0.02	0.24	0.71	0.69	0.05	0.29
MEND	0.96	<b>0.95</b>	0.96	0.96	0.31	0.31	0.27	0.29	0.15	0.14	0.03	0.09
ROME	0.90	0.89	0.99	0.93	0.77	0.73	0.52	0.66	0.31	0.28	0.03	0.14
MEMIT-M	0.84	0.81	0.99	0.88	0.73	0.72	0.95	0.79	0.70	0.65	0.82	0.72
DEFER	0.74	0.67	0.88	0.76	0.58	0.51	0.44	0.51	0.22	0.21	0.43	0.27
GRACE	0.97	0.41	0.98	0.73	<b>0.97</b>	0.2	<b>1.00</b>	0.58	<b>0.95</b>	0.08	<b>0.98</b>	0.42
WISE	0.97	<b>0.95</b>	0.98	0.97	0.79	0.73	0.91	0.80	0.59	0.57	0.92	0.68
<b>REPAIR</b>	<b>0.98</b>	<b>0.95</b>	<b>1.00</b>	<b>0.98</b> ↑	0.93	<b>0.90</b>	0.93	<b>0.92</b> ↑	0.81	<b>0.80</b>	0.92	<b>0.84</b> ↑
DeepSeek-R1-1.5B (WikiBigEdit)												
FT-L	0.71	0.68	0.93	0.77	0.26	0.20	0.76	0.34	0.13	0.11	0.37	0.17
FT-EWC	0.93	0.91	0.33	0.65	0.70	0.70	0.18	0.45	0.42	0.41	0.07	0.23
MEND	0.91	0.87	0.95	0.91	0.43	0.38	0.10	0.25	0.24	0.23	0.08	0.16
ROME	0.86	0.83	0.97	0.88	0.72	0.71	0.67	0.70	0.18	0.18	0.02	0.09
MEMIT-M	0.86	0.87	0.97	0.90	0.78	0.77	0.82	0.79	0.54	0.51	0.77	0.60
DEFER	0.68	0.58	0.47	0.35	0.63	0.61	0.51	0.58	0.17	0.15	0.33	0.20
GRACE	0.96	0.47	<b>0.99</b>	0.76	<b>0.93</b>	0.24	<b>0.91</b>	0.59	<b>0.76</b>	0.13	0.89	0.44
WISE	0.89	0.91	0.98	0.93	0.76	0.74	0.89	0.79	0.64	0.65	0.83	0.70
<b>REPAIR</b>	<b>0.98</b>	<b>0.93</b>	0.98	<b>0.96</b> ↑	0.84	<b>0.83</b>	<b>0.91</b>	<b>0.86</b> ↑	0.71	<b>0.69</b>	<b>0.90</b>	<b>0.76</b> ↑

Table 4: Main editing results for Hallucination task (SelfCheckGPT).

Method	$N = 1$		$N = 30$		$N = 120$		$N = 500$	
	Rel. ( $PPL \downarrow$ )	Loc. ( $\uparrow$ )	Rel. ( $\downarrow$ )	Loc. ( $\uparrow$ )	Rel. ( $\downarrow$ )	Loc. ( $\uparrow$ )	Rel. ( $\downarrow$ )	Loc. ( $\uparrow$ )
LLaMA-3-8B								
FT-L	4.27	0.96	3.15	0.71	34.52	0.43	51.31	0.26
FT-EWC	2.18	0.24	3.51	0.09	2.90	0.21	3.48	0.24
MEND	5.34	0.87	1.24	0.86	9.17	0.89	564.9	0.00
ROME	1.88	0.99	2.47	0.94	84.56	0.03	73.4	0.02
MEMIT-M	1.62	<b>1.00</b>	1.78	0.99	8.03	0.99	7.43	0.94
DEFER	<b>1.29</b>	0.23	4.12	0.28	8.91	0.19	15.16	0.12
GRACE	2.21	<b>1.00</b>	8.67	<b>1.00</b>	7.24	<b>1.00</b>	6.18	<b>1.00</b>
WISE	1.91	<b>1.00</b>	1.59	<b>1.00</b>	1.14	0.99	2.08	0.99
<b>REPAIR</b>	1.43	<b>1.00</b>	1.37	<b>1.00</b>	<b>1.12</b>	<b>1.00</b>	<b>1.91</b>	<b>1.00</b>

pruning/reassembly introduces transient instability before sufficient error signals accumulate; however, at  $N=1000$  their performance drops sharply, while REPAIR's dynamic adjustment preserves robustness and achieves the best overall metric. The error distribution can be seen in Appendix.

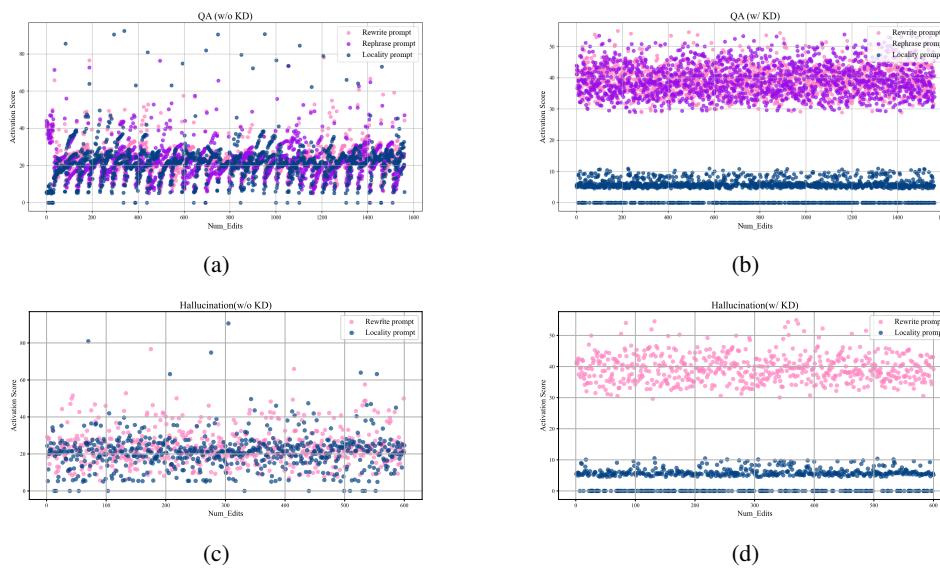


Figure 4: **Activation Score Visualization.** Results on LLaMA-3 for the WikiBigEdit dataset (N=1550) for the QA task and the SelfCheckGPT dataset for hallucination (N=600).

Figure 3 further addresses **Q1** regarding the effectiveness of distillation. For external memory-based editors, the ability to select the correct network for inference directly determines editing performance. The activation score, which serves as a critical routing criterion in memory networks, must exhibit statistically significant differences between new knowledge and irrelevant knowledge to ensure both reliability and locality of edits. As shown in Figure 4 (a) and (c), prior methods relying solely on triple-boundary loss fail to adequately separate the activation scores of *Dataedit*, *Datarephrase*, and *Dataloc*, particularly in large-scale continual editing scenarios, leading to a breakdown of the routing mechanism. This deficiency fundamentally limits their editing performance. In contrast, by introducing inner-batch knowledge distillation, sample filtering, and samples reintegration, KD, as shown in Figure 4 (b) and (d), achieves a clear separation among the three types of samples, thereby ensuring the proper functioning of the routing mechanism.

### 3.3 ABLATION STUDIES

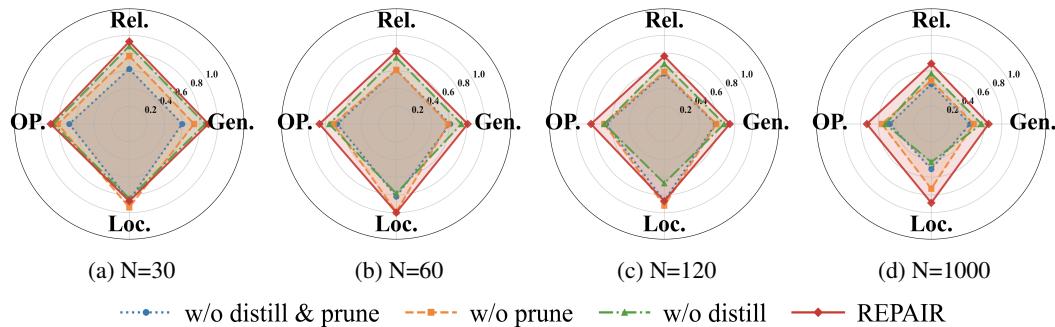


Figure 5: **Performance comparison of different components.** Each radar chart shows performance on four metrics: Rel., gen., loc., and OP. on Qwen2.5 with ZsRE.

The evaluation of the overhead and throughput of REPAIR can be found in the Appendix 7. To answer **Q6**, we conducted comprehensive evaluations across four dimensions to assess the effectiveness of each component of REPAIR and analyze critical hyperparameter sensitivity under different editing scales. Notably, REPAIR demonstrates robustness in large-scale editing scenarios

that prior methods fail to achieve. As the number of edits increases, REPAIR exhibits increasingly pronounced advantages in overall performance: effective routing ensures strong locality, while the error-feedback mechanism maintains continual reliability. As shown in Figure 5 (a)–(d), the relative contributions of REPAIR’s components vary across sample regimes but complement each other seamlessly. In small-scale edits, pruning with error feedback substantially improves reliability, while in large-scale scenarios, distribution-aware recognition and knowledge distillation become more critical. Regarding hyperparameter analysis in Figure 6, we observe distinct performance patterns: low thresholds fail to filter low-quality samples, limiting corrective opportunities; The total number of edits is limited, and the filtered erroneous samples cannot receive sufficient corrective training, which limits overall performance. A large number of erroneous samples in the early stage undergo continuous learning, causing the model to quickly fall into local optima, leading to catastrophic degradation of generalization. Subsequent learning yields minimal improvement, resulting in poor performance. In the upper-right quadrant, the absence of error feedback leaves many sub-optimal samples, and the model editing efficiency is relatively high, approximating an open-loop editing process. In the lower-right quadrant, the model training efficiency is the lowest, but excessive editing can introduce overfitting risks, wasting computational resources on edits with low marginal utility.

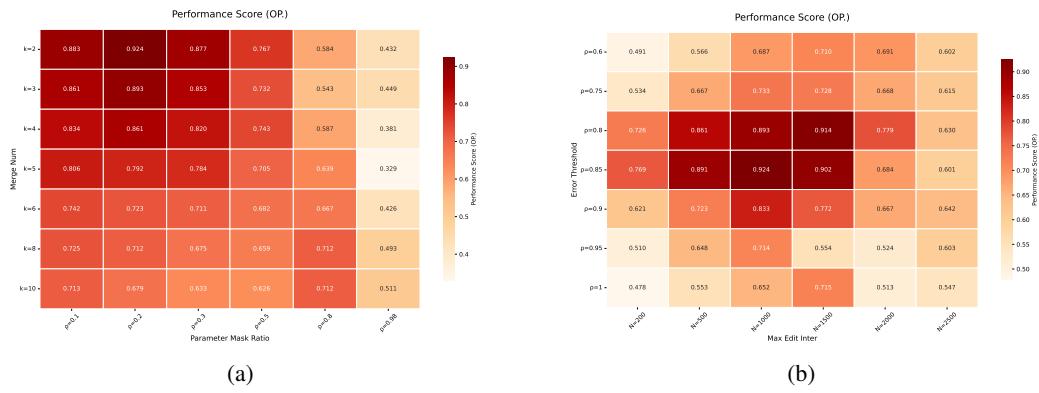


Figure 6: **Performance heatmap for the N=120 QA task on the LLaMA3 model.** Figure (a) shows the sensitivity analysis of two hyperparameters: the number of subspaces and the amount of updated parameters; Figure (b) analyzes the impact of error threshold and maximum iteration count on performance, with optimal performance observed at intermediate values.

## 4 CONCLUSION

In this work, we proposed **REPAIR**, a robust framework for lifelong model editing integrating error closed-loop feedback, inner-batch knowledge distillation, and loss-aware subspaces merging. Extensive experiments demonstrate that REPAIR maintains high performance under small-scale edits and exhibits remarkable robustness in large-scale editing scenarios, consistently outperforming existing baselines. These results highlight the potential of combining memory-aware strategies with optimization-driven editing for reliable and precise model updates. The intra-group distillation explicitly encourages feature alignment among similar samples, guiding the elimination and recombination of inconsistent samples. The loss-aware merging assigns higher weights to subspaces achieving lower training loss, effectively preserving reliable knowledge and reducing information dilution. Extensive experiments show that REPAIR consistently improves reliability and generalization, and demonstrates clear advantages in large-scale editing scenarios, highlighting the effectiveness of co-ordinated sample-level alignment and global reliability-aware merging.

## REFERENCES

- Nicola De Cao, Wilker Aziz, and Ivan Titov. Editing factual knowledge in language models. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 6491–6506, Online and Punta Cana, Dominican Republic, 2021. Association for Computational Linguistics.

- 486           tion for Computational Linguistics. doi: 10.18653/v1/2021.emnlp-main.522. URL <https://aclanthology.org/2021.emnlp-main.522/>.
- 487
- 488
- 489       Roi Cohen, Eden Biran, Ori Yoran, Amir Globerson, and Mor Geva. Evaluating the ripple effects  
490       of knowledge editing in language models. *Transactions of the Association for Computational  
491       Linguistics*, 12:283–298, 2024. doi: 10.1162/tacl\_a\_00644. URL <https://aclanthology.org/2024.tacl-1.16/>.
- 492
- 493       Mehrdad Farajtabar, Navid Azizan, Alex Mott, and Ang Li. Orthogonal gradient descent for con-  
494       tinual learning. In Silvia Chiappa and Roberto Calandra (eds.), *Proceedings of the Twenty Third  
495       International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 108 of *Pro-  
496       ceedings of Machine Learning Research*, pp. 3762–3773. PMLR, Aug 2020. URL <https://proceedings.mlr.press/v108/farajtabar20a.html>.
- 497
- 498
- 499       Enrico Fini, Victor G. Turrisi da Costa, Xavier Alameda-Pineda, Elisa Ricci, Karteek Alahari, and  
500       Julien Mairal. Self-supervised models are continual learners. In *Proceedings of the IEEE/CVF  
501       Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 9621–9630, June 2022.  
502       URL [https://openaccess.thecvf.com/content/CVPR2022/html/Fini\\_Self-Supervised\\_Models\\_Are\\_Continual\\_Learners\\_CVPR\\_2022\\_paper.html](https://openaccess.thecvf.com/content/CVPR2022/html/Fini_Self-Supervised_Models_Are_Continual_Learners_CVPR_2022_paper.html).
- 503
- 504
- 505       Ayush Gupta, Han Liu, Agastya Sharma, Di Jin, Neil Zhenqiang Gong, Sameer Singh, and Chenhao  
506       Tan. Rebuilding ROME: Resolving model collapse during sequential model editing. In *Proceed-  
507       ings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*,  
508       2024. URL <https://aclanthology.org/2024.emnlp-main.1210/>.
- 509
- 510       Thomas Hartvigsen, Swami Sankaranarayanan, Hamid Palangi, Yoon Kim, and Marzyeh  
511       Ghassemi. Aging with GRACE: Lifelong model editing with discrete key-value adap-  
512       tors. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL  
513       [https://proceedings.neurips.cc/paper\\_files/paper/2023/file/95b6e2ff961580e03c0a662a63a71812-Paper-Conference.pdf](https://proceedings.neurips.cc/paper_files/paper/2023/file/95b6e2ff961580e03c0a662a63a71812-Paper-Conference.pdf). NeurIPS  
514       2023.
- 515
- 516       James Kirkpatrick, Razvan Pascanu, Neil C. Rabinowitz, Joel Veness, Guillaume Desjardins, An-  
517       drei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwińska, Demis  
518       Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic for-  
519       getting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–  
520       3526, 2017. doi: 10.1073/pnas.1611835114. URL <https://www.pnas.org/doi/10.1073/pnas.1611835114>.
- 521
- 522       Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris  
523       Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion  
524       Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, , and Slav  
525       Petrov. Natural questions: A benchmark for question answering research. *Transactions of the  
526       Association for Computational Linguistics*, 7:452–466, 2019.
- 527
- 528       Omer Levy, Minjoon Seo, Eunsol Choi, and Luke Zettlemoyer. Zero-shot relation extraction via  
529       reading comprehension. In *Proceedings of the 21st Conference on Computational Natural Lan-  
530       guage Learning (CoNLL 2017)*, pp. 333–342, Vancouver, Canada, August 2017. Association for  
531       Computational Linguistics. doi: 10.18653/v1/K17-1034. URL <https://aclanthology.org/K17-1034/>.
- 532
- 533       Zhizhong Li and Derek Hoiem. Learning without forgetting. *IEEE Transactions on Pattern Analysis  
534       and Machine Intelligence*, 40(12):2935–2947, 2018. doi: 10.1109/TPAMI.2017.2773081. URL  
535       <https://doi.org/10.1109/TPAMI.2017.2773081>.
- 536
- 537       David Lopez-Paz and Marc’Aurelio Ranzato. Gradient episodic memory for contin-  
538       ual learning. In *Advances in Neural Information Processing Systems 30 (NeurIPS  
539       2017)*, pp. 6467–6476, 2017. URL <https://papers.nips.cc/paper/7225-gradient-episodic-memory-for-continual-learning>.

- 540 Arun Mallya and Svetlana Lazebnik. Packnet: Adding multiple tasks to a single net-  
 541 work by iterative pruning. In *Proceedings of the IEEE/CVF Conference on Computer Vi-  
 542 sion and Pattern Recognition (CVPR)*, pp. 7765–7773, 2018. doi: 10.1109/CVPR.2018.  
 543 00810. URL [https://openaccess.thecvf.com/content\\_cvpr\\_2018/papers/Mallya\\_PackNet\\_Adding\\_Multiple\\_CVPR\\_2018\\_paper.pdf](https://openaccess.thecvf.com/content_cvpr_2018/papers/Mallya_PackNet_Adding_Multiple_CVPR_2018_paper.pdf).
- 544 Michael McCloskey and Neal J. Cohen. Catastrophic interference in connectionist networks: The  
 545 sequential learning problem. In *The Psychology of Learning and Motivation*, volume 24, pp.  
 546 109–165. Academic Press, 1989. doi: 10.1016/S0079-7421(08)60536-8. URL <https://www.sciencedirect.com/science/article/pii/S0079742108605368>.
- 547 Sanket Vaibhav Mehta, Darshan Patil, Sarath Chandar, and Emma Strubell. An empirical investi-  
 548 gation of the role of pre-training in lifelong learning. *Journal of Machine Learning Research*, 24  
 549 (214):1–50, 2023. URL <https://www.jmlr.org/papers/v24/22-0496.html>.
- 550 Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual  
 551 associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022a.  
 552 URL [https://proceedings.neurips.cc/paper\\_files/paper/2022/file/6f1d43d5a82a37e89b0665b33bf3a182-Paper-Conference.pdf](https://proceedings.neurips.cc/paper_files/paper/2022/file/6f1d43d5a82a37e89b0665b33bf3a182-Paper-Conference.pdf). NeurIPS 2022.
- 553 Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual  
 554 associations in gpt. *Advances in neural information processing systems*, 35:17359–17372, 2022b.
- 555 Kevin Meng, David Bau, Alex Andonian, Yonatan Belinkov, and David Bau Lab. Counterfact:  
 556 A benchmark for evaluating knowledge editing locality and generalization. <https://rome.baulab.info/>, 2022c. Dataset introduced alongside ROME.
- 557 Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, and David Bau. Mass-editing  
 558 memory in a transformer. In *International Conference on Learning Representations (ICLR)*, 2023.  
 559 URL <https://openreview.net/forum?id=MkbcAHIYgyS>.
- 560 Eric Mitchell, Charles Lin, Antoine Bosselut, Christopher D. Manning, and Chelsea Finn. Fast  
 561 model editing at scale. In *International Conference on Learning Representations (ICLR)*, 2022a.  
 562 URL <https://openreview.net/pdf?id=0DcZxeWfOPT>.
- 563 Eric Mitchell, Charles Lin, Antoine Bosselut, Christopher D. Manning, and Chelsea Finn.  
 564 Memory-based model editing at scale. In *Proceedings of the 39th International Conference  
 565 on Machine Learning (ICML)*, volume 162 of *Proceedings of Machine Learning Research*,  
 566 pp. 15828–15846. PMLR, 2022b. URL <https://proceedings.mlr.press/v162/mitchell22a.html>.
- 567 Xuming Ran, Juntao Yao, Yusong Wang, Mingkun Xu, and Dianbo Liu. Brain-inspired continual  
 568 pre-trained learner via silent synaptic consolidation. *ArXiv*, abs/2410.05899, 2024. URL <https://api.semanticscholar.org/CorpusID:273228983>.
- 569 Andrei A. Rusu, Neil C. Rabinowitz, Guillaume Desjardins, Hubert Soyer, James Kirkpatrick, Koray  
 570 Kavukcuoglu, Razvan Pascanu, and Raia Hadsell. Progressive neural networks. *arXiv preprint  
 571 arXiv:1606.04671*, 2016. URL <https://arxiv.org/abs/1606.04671>.
- 572 Hanul Shin, Jung Kwon Lee, Jaehong Kim, and Jiwon Kim. Continual learning with  
 573 deep generative replay. In *Advances in Neural Information Processing Systems 30  
 574 (NeurIPS 2017)*, pp. 2990–2999, 2017. URL <https://papers.nips.cc/paper/6892-continual-learning-with-deep-generative-replay>.
- 575 Peng Wang, Zexi Li, Ningyu Zhang, Ziwen Xu, Yunzhi Yao, Yong Jiang, Pengjun Xie, Fei Huang,  
 576 and Huajun Chen. Wise: Rethinking the knowledge memory for lifelong model editing of large  
 577 language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024a. doi:  
 578 10.48550/arXiv.2405.14768. URL <https://arxiv.org/abs/2405.14768>. NeurIPS  
 579 2024. Also available as arXiv:2405.14768.
- 580 Song Wang, Yaochen Zhu, Haochen Liu, Zaiyi Zheng, Chen Chen, and Jundong Li. Knowl-  
 581 edge editing for large language models: A survey. *ACM Computing Surveys*, 2024b. doi:  
 582 10.1145/3698590. URL <https://dl.acm.org/doi/10.1145/3698590>. Also available  
 583 as arXiv:2310.16218.

594 Zhenyi Wang, Zhen Zhang, Xiaojuan E, Ziyu Zhang, Zhipeng Luo, Zhipeng He, and Guangyao Li.  
 595 A comprehensive survey on continual learning: From definitions to applications. *arXiv preprint*  
 596 *arXiv:2302.00487*, 2023. URL <https://arxiv.org/abs/2302.00487>.

597  
 598 Maurice Weber, Daniel Y. Fu, Quentin Anthony, Yonatan Oren, Shane Adams, Anton Alexandrov,  
 599 Xiaozhong Lyu, Huu Nguyen, Xiaozhe Yao, Virginia Adams, Ben Athiwaratkun, Rahul Chal-  
 600 lamala, Kezhen Chen, Max Ryabinin, Tri Dao, Percy Liang, Christopher Ré, Irina Rish, and  
 601 Ce Zhang. Redpajama: an open dataset for training large language models. In Amir Globersons,  
 602 Lester Mackey, Danielle Belgrave, Angela Fan, Ulrich Paquet, Jakub M. Tomczak, and Cheng  
 603 Zhang (eds.), *Advances in Neural Information Processing Systems 38: Annual Conference on*  
 604 *Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, De-*  
 605 *cember 10 - 15, 2024, 2024*. URL [http://papers.nips.cc/paper\\_files/paper/2024/hash/d34497330b1fd6530f7afdf86d0df9f76-Abstract-Datasets\\_and\\_Benchmarks\\_Track.html](http://papers.nips.cc/paper_files/paper/2024/hash/d34497330b1fd6530f7afdf86d0df9f76-Abstract-Datasets_and_Benchmarks_Track.html).

606  
 607 Prateek Yadav, Derek Tam, Leshem Choshen, Colin Raffel, and Mohit Bansal. Ties-merging: Re-  
 608 solving interference when merging models, 2023. URL <https://arxiv.org/abs/2306.01708>.

609  
 610 Ningyu Zhang, Bozhong Tian, Siyuan Cheng, Xiaozhuan Liang, Yi Hu, Kouying Xue, Yanjie Gou,  
 611 Xi Chen, and Huajun Chen. Instructedit: Instruction-based knowledge editing for large lan-  
 612 guage models. In *Proceedings of the 33rd International Joint Conference on Artificial Intelli-  
 613 gence (IJCAI 2024)*, 2024. doi: 10.24963/ijcai.2024/733. URL <https://www.ijcai.org/proceedings/2024/0733.pdf>.

## 614 A STATEMENT

### 615 A.1 ETHICS STATEMENT

616 This work studies safe, auditable editing of large language models using only publicly available  
 617 datasets (ZsRE, WikiBigEdit, and a hallucination set) and off-the-shelf pretrained models; no human  
 618 subjects or personally identifiable data were collected. We follow all dataset/model licenses and the  
 619 double-blind review policy. Potential risks include misuse of editing to inject misinformation or to  
 620 weaken safety constraints, and unintended spillover of edits to unrelated behaviors. To mitigate these  
 621 risks, our framework emphasizes locality and closed-loop error checks before and after integration,  
 622 and we report reliability-generalization-and locality metrics to surface side effects. Upon release,  
 623 we will include guardrails such as edit logs, validation suites, reversible edits, and instructions for  
 624 responsible use. These design choices align with REPAIR’s stated goal of precise updates with  
 625 locality safeguards.

### 626 A.2 REPRODUCIBILITY STATEMENT

627 We will release our code, configs, and seeds to reproduce all results end-to-end after acceptance.  
 628 Scripts fetch data/models, fix environments, and regenerate all tables/figures with the exact metrics  
 629 (Rel./Gen./Loc./OP., PPL); hardware and hyperparameters are documented.

### 630 A.3 AI USAGE STATEMENT

631 We used large language model-based tools during writing and implementation for text polishing,  
 632 grammar and usage checks, and programming assistance (e.g., example code, refactoring, com-  
 633 ments, and script templates). All AI-generated suggestions were reviewed, revised, and validated  
 634 by the authors. The experimental design, data processing, result analysis, and conclusions were  
 635 conducted independently by the authors; AI tools do not constitute authorship or academic credit.  
 636 No sensitive or restricted data were provided to the tools, and they were not used to automatically  
 637 generate experimental results or to replace essential human judgment.

648      **B RELATED WORK**

649

650      **B.1 CONTINUAL LEARNING**

651

653      Continual Learning (CL)—also known as Incremental Learning or Lifelong Learning—aims to enable models to learn sequentially from a stream of tasks without forgetting previously acquired knowledge. The core challenge in CL is catastrophic forgetting, where adapting to new tasks leads to a significant degradation in performance on earlier tasks Kirkpatrick et al. (2017); McCloskey & Cohen (1989). To address this, numerous methods have been proposed, which can be broadly categorized into five groups: regularization-based, replay-based, optimization-based, representation-based, and architecture-based approaches.

654

656      Regularization-based methods mitigate forgetting by adding constraints to the loss function to preserve important parameters or behaviors from previous tasks. For example, Elastic Weight Consolidation (EWC) leverages Fisher information to regularize parameter updates Kirkpatrick et al. (2017), while Learning without Forgetting (LwF) uses knowledge distillation to maintain output consistency Li & Hoiem (2018).

657

659      Replay-based methods retain or generate samples from previous tasks to approximate old data distributions. Experience replay stores a subset of prior samples in a memory buffer Lopez-Paz & Ranzato (2017), whereas generative replay synthesizes pseudo-samples using deep generative models such as GANs or VAEs Shin et al. (2017).

660

662      Optimization-based methods manipulate the optimization process itself to avoid interference between tasks. Gradient Episodic Memory (GEM) projects gradients so as not to increase loss on previous tasks Lopez-Paz & Ranzato (2017), while Orthogonal Gradient Descent (OGD) promotes updates that are orthogonal to gradient directions associated with past tasks Farajtabar et al. (2020).

663

665      Representation-based methods focus on learning robust and transferable features that are less prone to forgetting. Self-supervised learning Fini et al. (2022) and large-scale pre-training Mehta et al. (2023) have been shown to bolster CL performance by providing more stable representations.

666

668      Architecture-based methods Ran et al. (2024); Rusu et al. (2016); Mallya & Lazebnik (2018) dynamically expand or partition the network to allocate task-specific parameters. Progressive Networks add new columns for each incoming task with lateral connections to prior columns Rusu et al. (2016), while PackNet iteratively prunes and reuses weights to free capacity for new tasks Mallya & Lazebnik (2018).

669

671      Recent trends extend CL to more realistic and challenging settings, including class-incremental learning (CIL), task-free CL (TFCL), online CL (OCL), and applications across object detection, semantic segmentation, reinforcement learning, and natural language processing Wang et al. (2023).

672

674      **B.2 MODEL EDITING**

675

677      Model editing targets post-hoc modification of a trained model’s behavior to insert, correct, or remove specific knowledge, ideally without harming unrelated capabilities. A common taxonomy distinguishes (i) *direct/training-free* parameter edits, (ii) *learning-based* editors that predict weight updates, and (iii) *semi-parametric* systems that externalize edits via retrieval or memory; recent surveys consolidate definitions, benchmarks, and open challenges Wang et al. (2024b).

678

679      ROME locates causal mediators of factual associations in mid-layer feed-forward (MLP) modules of Transformers and applies a rank-one update to edit a single fact Meng et al. (2022a). MEMIT extends this idea to *mass editing*, deriving multi-layer closed-form updates that scale to thousands of edits in large models while maintaining stronger locality than prior methods Meng et al. (2023). Although effective, subsequent analyses highlight stability issues under *sequential* edits and propose remedies Gupta et al. (2024).

680

681      Early work framed editing as learning a small hypernetwork to predict weight deltas from an edit specification: KnowledgeEditor (KE) learns constrained updates to change a model’s factual prediction while preserving behavior on paraphrases Cao et al. (2021). MEND trains lightweight editor networks to transform fine-tuning gradients, enabling fast, local edits at scale across architectures

682

702  
703  
704  
705  
706  
707  
708  
709  
710  
711  
712  
713  
714  
715  
716  
717  
718  
719  
720  
721  
722  
723  
724  
725  
726  
727  
728  
729  
730  
731  
732  
733  
734  
735  
736  
737  
738  
739  
740  
741  
742  
743  
744  
745  
746  
747  
748  
749  
750  
751  
752  
753  
754  
755  
Table 5: Dataset statistics

Task	Editing Data	N	Pre-edit(LLaMA/Qwen)	Locality Data	
QA	ZsRE	1000	0.25/0.21 ACC	NQ	Kwiatkowski et al. (2019)
	WikiBigEdit	500K	0.36/0.32 ACC	NQ	
Hallu.	SelfCheckGPT	600	28.7/29.1 PPL	RedPajama	Weber et al. (2024)

Table 6: Hyperparameter settings

ZsRE on LLaMA-3						
	HYPER	VALUE	HYPER	VALUE	HYPER	VALUE
Mask ratio	0.20	Edit_lr	0.90	Err.Thresh	0.85	
$\lambda_a$	1.00	$\lambda_{KD}$	1.00	Max_iter	10000	
Temperature	2.00	Act ratio	0.20	Layer_ID	29.00	
$\gamma_1$	2.00	$\gamma_2$	20.00	$\gamma$	10.00	
$n_{iter}$	30.00	$\lambda$	0.20	Act_ratio	0.30	
ZsRE on Qwen2.5						
Mask ratio	0.20	Edit_lr	0.90	Err.Thresh	0.85	
$\lambda_a$	2.00	$\lambda_{KD}$	1.00	Max_iter	10000	
Temperature	2.00	Act ratio	0.88	Layer_ID	23.00	
$\gamma_1$	5.00	$\gamma_2$	20.00	$\gamma$	10.0	
$n_{iter}$	50.00	$\lambda$	0.30	Act_ratio	0.30	
Selfcheck GPT on LLaMA-3-8B						
Mask ratio	0.20	Edit_lr	1.00	Err.Thresh	0.85	
$\lambda_a$	5.00	$\lambda_{KD}$	1.00	Max_iter	5000	
Temperature	2.00	Act ratio	0.88	Layer_ID	27.00	
$\gamma_1$	5.00	$\gamma_2$	20.00	$\gamma$	10.00	
$n_{iter}$	50.00	$\lambda$	0.20	Act_ratio	0.80	

Mitchell et al. (2022a). Instruction-driven variants further condition edits on natural-language instructions to improve usability and control Zhang et al. (2024).

Semi-parametric approaches such as SERAC store edits in an external key–value memory and learn to route between the base model and retrieved counterfactuals, achieving strong reliability and specificity without permanently altering base parameters Mitchell et al. (2022b). This design is attractive when edits must be audited, reverted, or scoped to contexts.

Editing methods are typically assessed along *reliability* (does the change take effect), *locality/specificty* (does unrelated behavior remain intact), and *generalization* (do edits transfer to paraphrases and contexts). Standard benchmarks include CounterFact and zsRE Meng et al. (2022c); Levy et al. (2017). Recent studies examine *ripple effects* beyond targeted facts, revealing broader side impacts on reasoning and distributed knowledge, and call for more rigorous, stress-testing evaluations Cohen et al. (2024). Overall, direct, learning-based, and semi-parametric approaches offer complementary trade-offs in edit scalability, controllability, and safety; combining precise localization with guardrails (e.g., retrieval gating, edit scopes, or validation filters) remains an active direction Wang et al. (2024b).

## C EXPERIMENTS DETAILS

The experiment details are given in Table 5, and hyperparameters are in Table 6.

Under identical hardware and batch configurations, the WISE baseline exhibits lower per-unit overhead. REPAIR demonstrates a similar scaling slope but with a higher intercept, primarily attributable to:

Table 7: Main results for QA on DeepSeek-R1-1.5B  $N$ : Num Edits.

Method	$N = 1$			$N = 30$			$N = 120$			$N = 1000$		
	Rel.	Gen.	Loc.	OP.	Rel.	Gen.	Loc.	OP.	Rel.	Gen.	Loc.	OP.
DeepSeek-R1-1.5B (ZsRE)												
FT-L	0.43	0.42	0.95	0.56	0.32	0.33	0.46	0.36	0.21	0.21	0.15	0.19
FT-EWC	0.97	<b>0.94</b>	0.15	0.52	0.82	0.81	0.02	0.24	0.63	0.64	0.02	0.20
MEND	0.95	<b>0.94</b>	0.98	0.96	0.42	0.42	0.18	0.32	0.18	0.12	0.07	0.11
ROME	0.87	0.87	0.99	0.91	0.66	0.64	0.72	0.67	0.17	0.18	0.09	0.14
MEMIT-M	0.88	0.87	0.99	0.91	0.71	0.72	0.92	0.78	0.63	0.65	0.78	0.68
DEFER	0.62	0.60	0.82	0.67	0.58	0.57	0.57	0.57	0.34	0.31	0.23	0.29
GRACE	<b>0.98</b>	0.31	0.99	0.67	<b>0.92</b>	0.22	<b>0.98</b>	0.58	<b>0.89</b>	0.13	<b>1.00</b>	0.49
WISE	0.92	0.90	<b>1.00</b>	0.94	0.86	0.85	0.92	0.88	0.72	0.72	0.87	<b>0.77</b>
<b>REPAIR</b>	0.93	0.93	<b>1.00</b>	<b>0.95</b>	0.91	<b>0.89</b>	0.87	<b>0.89↑</b>	0.74	<b>0.74</b>	0.82	<b>0.77↑</b>

Table 8: Main results for QA (ZeRE) on multi-model editing with error distribution.

Method	$N = 1$			$N = 30$		
	Rel.	Gen.	Loc.	Rel.	Gen.	Loc.
LLaMA-3-8B	$0.94 \pm 0.008$	$0.92 \pm 0.01$	$1.00^{+0.00}_{-0.02}$	$0.93 \pm 0.003$	$0.90 \pm 0.003$	$0.87 \pm 0.004$
Qwen2.5-7B	$0.98 \pm 0.02$	$0.95 \pm 0.03$	$1.00^{+0.00}_{-0.02}$	$0.93 \pm 0.04$	$0.90 \pm 0.03$	$0.93 \pm 0.01$
DeepSeek-R1	$0.93 \pm 0.02$	$0.92 \pm 0.03$	$0.99 \pm 0.01$	$0.91 \pm 0.01$	$0.89 \pm 0.03$	$0.87 \pm 0.01$
GPT2-XL	$0.91 \pm 0.03$	$0.92 \pm 0.03$	$0.99 \pm 0.01$	$0.88 \pm 0.03$	$0.88 \pm 0.02$	$0.84 \pm 0.01$

Method	$N = 120$			$N = 1000$		
	Rel.	Gen.	Loc.	Rel.	Gen.	Loc.
LLaMA-3-8B	$0.76 \pm 0.03$	$0.74 \pm 0.02$	$1.00^{+0.00}_{-0.04}$	$0.68 \pm 0.05$	$0.65 \pm 0.01$	$0.89 \pm 0.04$
Qwen2.5-7B	$0.81 \pm 0.04$	$0.80 \pm 0.05$	$0.92 \pm 0.03$	$0.72 \pm 0.05$	$0.70 \pm 0.04$	$0.67 \pm 0.03$
DeepSeek-R1	$0.74 \pm 0.03$	$0.74 \pm 0.04$	$0.82 \pm 0.05$	$0.59 \pm 0.02$	$0.57 \pm 0.01$	$0.61 \pm 0.03$
GPT2-XL	$0.79 \pm 0.02$	$0.77 \pm 0.01$	$0.80 \pm 0.03$	$0.61 \pm 0.03$	$0.62 \pm 0.01$	$0.68 \pm 0.02$

- Distribution-aware clustering and reorganization;
- Additional forward/backward passes for in-batch distillation;
- The triggering frequency and cost of closed-loop pruning and retraining at large scales;
- The final **merging (TIES)** cost.

As  $N$  increases, the runtime curves of ROME, MEND, and FT-L exhibit significantly steeper growth, becoming substantially expensive or nearly infeasible at  $N = 10^3$ .

**Throughput** demonstrates that WISE maintains approximately  $\sim 1.8$  edits/min at large  $N$ , followed by GRACE. REPAIR achieves  $\sim 0.8\text{-}0.9$  edits/min at scale, lower than WISE and MEMIT-M, consistent with expectations given its additional computational procedures.

Error bars (representing standard deviation across multiple runs) indicate that REPAIR exhibits slightly higher variance than WISE, attributable to fluctuations in retriggering frequency and sample distribution characteristics.

**Relative overhead** shows the time ratio of REPAIR to WISE increasing from  $\sim 1.6\times$  to  $\sim 2.2\times$  with increasing scale.

## D THEORETICAL ANALYSIS AND PROOF SKETCHES

We now provide theoretical justifications for the stability and convergence of the proposed REPAIR framework. We introduce formal assumptions and derive lemmas and theorems that characterize the behavior of our method.

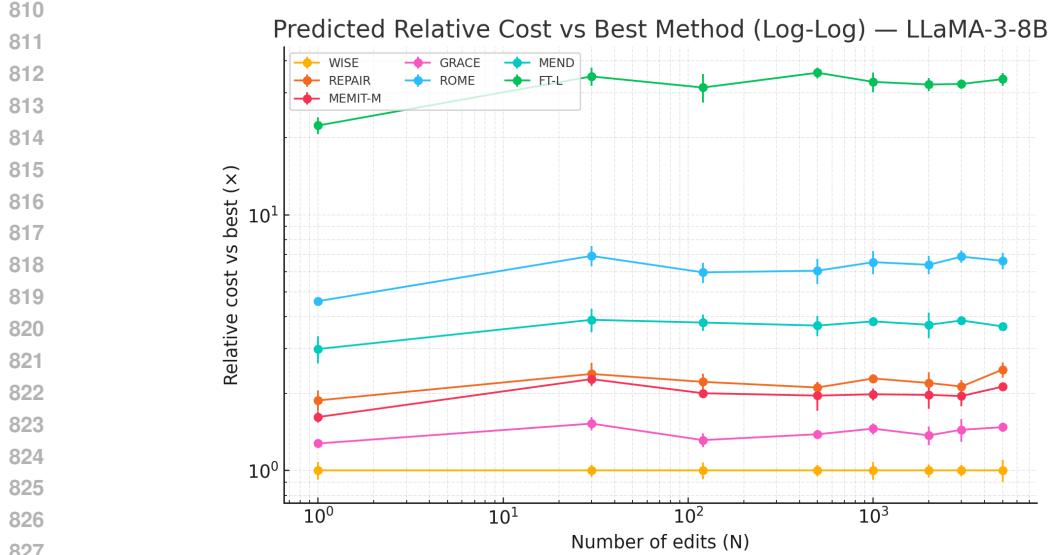


Figure 7: **Cost-Performance Assessment.** The total runtime of each method scales approximately linearly with the editing scale  $N$ , appearing as straight lines with slopes close to 1 in log-log coordinates. This indicates that the primary overhead is proportional to the number of edited entries.

## D.1 PRELIMINARIES

**Assumption 1** (Standard Optimization Setting). *We assume that the loss function  $\mathcal{L}(\Theta)$  is  $L$ -smooth, i.e.,*

$$\|\nabla \mathcal{L}(\Theta_1) - \nabla \mathcal{L}(\Theta_2)\| \leq L\|\Theta_1 - \Theta_2\|,$$

*and bounded below by  $\mathcal{L}^* > -\infty$ . Learning rates satisfy  $\eta_t > 0$  and  $\sum_t \eta_t = \infty$ ,  $\sum_t \eta_t^2 < \infty$ .*

## D.2 STABILITY OF MASKED GRADIENT UPDATES

**Lemma 1** (Norm Bound under Masked Updates). *Let  $g_i = \nabla_{W'_{v,i}} \mathcal{L}$  and  $M_i$  be a Bernoulli mask. Then the masked update*

$$\Delta W'_{v,i} = -\eta(M_i \odot g_i)$$

*satisfies  $\|\Delta W'_{v,i}\|_2 \leq \eta \|g_i\|_2$ .*

*Proof.* Since  $M_i$  is a coordinate projection,  $M_i \odot g_i$  removes certain entries of  $g_i$  and never increases its magnitude. Hence  $\|M_i \odot g_i\|_2 \leq \|g_i\|_2$ . Multiplying by  $\eta$  yields the claim.  $\square$

**Theorem 1** (Inter-Shard Stability). *Assume masks  $\{M_i\}$  are sampled independently with overlap probability  $\rho^2$ . Then in expectation,*

$$\mathbb{E}[\langle M_i \odot g_i, M_j \odot g_j \rangle] = \rho^2 \langle g_i, g_j \rangle.$$

*Thus, masking reduces the expected conflict between gradients of different shards.*

*Proof.* For each coordinate  $p$ ,  $\Pr[M_i(p) = 1, M_j(p) = 1] = \rho^2$ . Therefore, the expected inner product between masked gradients is  $\rho^2$  times the original inner product. This reduces cross-shard interference and improves stability.  $\square$

## D.3 CLOSED-LOOP RE-TRIGGER ANALYSIS

**Assumption 2** (Error Reduction per Re-trigger). *Suppose that each re-trigger reduces the error rate of shard  $i$  by at least a fixed constant  $\delta > 0$ , unless it is already below the pruning threshold  $\tau_{\text{prune}}$ .*

864    **Lemma 2** (Linear Error Decrease). Let  $r_i^{(n)}$  denote the error rate after  $n$  re-triggers. Under Assumption 2,

$$r_i^{(n)} \leq r_i^{(0)} - n\delta.$$

868    **Theorem 2** (Finite-Time Convergence). If  $r_i^{(0)}$  is the initial error rate, then after at most

$$869 \quad N \geq \frac{r_i^{(0)} - \tau_{\text{prune}}}{\delta}$$

872    re-triggers, the error rate satisfies  $r_i^{(N)} \leq \tau_{\text{prune}}$ .

874    *Proof.* By Lemma 3,  $r_i^{(N)} \leq r_i^{(0)} - N\delta$ . Choosing  $N$  such that  $r_i^{(0)} - N\delta \leq \tau_{\text{prune}}$  ensures convergence below threshold in finite time.  $\square$

#### 877    D.4 OVERALL CONVERGENCE INTUITION

879    **Theorem 3** (Closed-Loop Stability of REPAIR). Under Assumptions 1 and 2, the iterative process  
880    combining masked updates, inner-batch distillation, and closed-loop re-trigger forms a contractive  
881    mapping in expectation. Consequently, the system converges to a stable edited state with a bounded  
882    error rate and without catastrophic forgetting.

883    *Proof Sketch.* Masked updates reduce the variance of parameter updates, inner-batch distillation  
884    aligns outputs across samples, and re-trigger guarantees finite-time reduction of shard-level error  
885    rates. Together, these components yield monotone improvement. By standard stochastic contraction  
886    arguments, the process converges to a fixed point characterized by consistent batch predictions and  
887    an error rate below  $\tau_{\text{prune}}$ .  $\square$

889    **Lemma 3** (Zero-variance at any global minimizer). Let  $\mu = \frac{1}{m} \sum_{i=1}^m o_i$  and  $\mathcal{L}_{\text{var}} = \frac{1}{m} \sum_i \|o_i - \mu\|^2$ . If not all  $o_i$  are equal, then  $\mathcal{L}_{\text{var}} > 0$ , while if  $o_1 = \dots = o_m = v$  (with  $\|v\| = 1$ ) then  $\mathcal{L}_{\text{var}} = 0$ . Hence every global minimizer of  $\mathcal{L}_{\text{KD}}$  on  $(\mathbb{S}^{d-1})^m$  must satisfy  $o_1 = \dots = o_m =: v$ .

892    **Lemma 4** (Unique global minimizer). Under the conclusion of Lemma 3, minimizing  $\mathcal{L}_{\text{KD}}(v) = \lambda(1 - \langle v, u \rangle)$  over  $\|v\| = 1$  gives the unique solution  $v^* = u$ . Therefore the unique global minimizer  
893    of  $\mathcal{L}_{\text{KD}}$  on  $(\mathbb{S}^{d-1})^m$  is  $S^* = [u, \dots, u]$ .

894    **Lemma 5** (Riemannian smoothness). Let  $\mathcal{M} = (\mathbb{S}^{d-1})^m$  and endow each block with the canonical  
895    metric. Then  $\mathcal{L}_{\text{KD}}$  is  $L_R$ -smooth on  $\mathcal{M}$  in the Riemannian sense: there exists a constant

$$896 \quad L_R \leq \frac{2\lambda}{m} + \frac{4\vartheta}{m}$$

900    such that for all  $S, S' \in \mathcal{M}$ ,  $\|\text{grad } \mathcal{L}_{\text{KD}}(S) - \text{grad } \mathcal{L}_{\text{KD}}(S')\| \leq L_R \text{dist}_{\mathcal{M}}(S, S')$ . Sketch. For  
901    each block  $o_i$ ,  $\nabla_{o_i} \mathcal{L}_{\text{cos}} = -(\lambda/m)u$  (constant), and  $\nabla_{o_i} \mathcal{L}_{\text{var}} = (2\vartheta/m)(o_i - \mu)$  with  $\mu$  depending  
902    linearly on  $\{o_j\}$ . Projecting to the tangent space by  $(I - o_i o_i^\top)$  and using the Lipschitzness of the  
903    projection map on  $\mathbb{S}^{d-1}$  yields the bound.

904    **Theorem 4** (Convergence of cosine+variance KD on the sphere). Consider Riemannian gradient  
905    descent on  $\mathcal{M} = (\mathbb{S}^{d-1})^m$ :

$$906 \quad o_i^{(t+1)} = R_{o_i^{(t)}}(-\eta_t \text{grad}_{o_i} \mathcal{L}_{\text{KD}}(S_t)) \quad (i = 1, \dots, m),$$

908    with the retraction  $R_o(v) = (o + v)/\|o + v\|$ . If the step sizes satisfy either (a) a constant stepsize  
909     $0 < \eta_t < 2/L_R$ , or (b) diminishing stepsizes  $\sum_t \eta_t = \infty$ ,  $\sum_t \eta_t^2 < \infty$ , then:

$$910 \quad \mathcal{L}_{\text{KD}}(S_t) \downarrow \mathcal{L}_{\text{KD}}(S^*), \quad \|\text{grad } \mathcal{L}_{\text{KD}}(S_t)\| \rightarrow 0,$$

912    and every limit point of  $\{S_t\}$  is a Riemannian critical point. By Lemma 4, the unique global mini-  
913    mizer is  $S^* = [u, \dots, u]$ ; thus the sequence converges to  $S^*$ .

915    *Proof sketch.* Riemannian smoothness (Lemma 5) on the compact manifold  $\mathcal{M}$  ensures the standard  
916    descent lemma and monotone decrease for RGD under  $0 < \eta < 2/L_R$ , implying convergence of  
917    function values and gradients to zero. By Lemmas 3–4, the only global minimizer is  $S^*$ , hence all  
918    limit points coincide with  $S^*$ .  $\square$

918 D.5 STABILITY OF MASKED GRADIENT UPDATES  
919920 Let  $g_i = \nabla_{W'_{v,i}} \mathcal{L} \in \mathbb{R}^d$ . A coordinate mask  $M_i \in \{0, 1\}^d$  acts by  $(M_i \odot g_i)_p = M_i(p) g_{i,p}$ .921 **Lemma 6** (Norm Bound under Masked Updates). *For any stepsize  $\eta > 0$ , the masked update  
922  $\Delta W'_{v,i} = -\eta(M_i \odot g_i)$  satisfies*

923 
$$\|\Delta W'_{v,i}\|_2 \leq \eta \|g_i\|_2.$$
  
924

925 *Proof.* Coordinate-wise,  $|M_i(p) g_{i,p}| \leq |g_{i,p}|$  because  $M_i(p) \in \{0, 1\}$ . Hence  $\|M_i \odot g_i\|_2 \leq \|g_i\|_2$ ,  
926 and multiplying by  $\eta$  yields the claim.  $\square$   
927928 **Theorem 5** (Inter-Shard Inner-Product Scaling). *Suppose that for each coordinate  $p$ , the masks  
929  $M_i(p), M_j(p) \in \{0, 1\}$  are sampled independently with*

930 
$$\Pr[M_i(p) = 1] = \Pr[M_j(p) = 1] = \rho, \quad 0 \leq \rho \leq 1,$$
  
931

932 *and masks are independent across coordinates and independent of  $g_i, g_j$ . Then, conditional on  
933  $g_i, g_j$ ,*

934 
$$\mathbb{E}[\langle M_i \odot g_i, M_j \odot g_j \rangle | g_i, g_j] = \rho^2 \langle g_i, g_j \rangle.$$
  
935

936 *In particular, masking scales the expected cross-shard alignment/conflict by the factor  $\rho^2$ .*  
937938 *Proof.* By linearity of expectation and independence, for each coordinate  $p$ ,  $\mathbb{E}[M_i(p)M_j(p)] =$   
939  $\mathbb{E}[M_i(p)]\mathbb{E}[M_j(p)] = \rho^2$ . Summing over  $p$  yields the result.  $\square$   
940

## D.6 CLOSED-LOOP RE-TRIGGER ANALYSIS

941 **Assumption 3** (Error Reduction per Re-trigger). *Let  $r_i^{(n)}$  denote the error rate of shard  $i$  after  
942  $n$  re-triggers. There exists  $\delta > 0$  such that each re-trigger reduces error by at least  $\delta$  whenever  
943  $r_i^{(n)} > \tau_{\text{prune}}$ .*944 **Lemma 7** (Piecewise-Linear Error Decrease). *Under Assumption 3, for all  $n \geq 0$ ,*

945 
$$r_i^{(n)} \leq \max\{\tau_{\text{prune}}, r_i^{(0)} - n\delta\}.$$
  
946

947 *Proof.* If  $r_i^{(k)} > \tau_{\text{prune}}$ , then  $r_i^{(k+1)} \leq r_i^{(k)} - \delta$ . Once  $r_i^{(k)} \leq \tau_{\text{prune}}$ , the bound  $r_i^{(n)} \leq \tau_{\text{prune}}$   
948 propagates for all  $n \geq k$ . Unrolling gives the stated maximum form.  $\square$   
949950 **Theorem 6** (Finite-Time Hitting the Pruning Threshold). *Let*  
951

952 
$$N_* = \left\lceil \frac{(r_i^{(0)} - \tau_{\text{prune}})_+}{\delta} \right\rceil \quad \text{where } (x)_+ := \max\{x, 0\}.$$
  
953

954 *After at most  $N_*$  re-triggers, we have  $r_i^{(N_*)} \leq \tau_{\text{prune}}$ .*  
955956 *Proof.* By Lemma 7, choose the smallest integer  $N_*$  such that  $r_i^{(0)} - N_*\delta \leq \tau_{\text{prune}}$ . Then  $r_i^{(N_*)} \leq$   
957  $\tau_{\text{prune}}$ .  $\square$   
958959 E ALGORITHMS  
960961 The pseudocode for error feedback, network pruning, sample knowledge distillation and reintegration,  
962 and the loss-based weighted ties merge strategy is as follows:  
963964  
965  
966  
967  
968  
969  
970  
971

---

972  
973  
974  
975  
976

**Algorithm 1** REPAIR: Closed-Loop Lifelong Model Editing (Training)

---

977 **Require:** Pretrained model  $f_{\theta_0}$ ; target FFN value matrix  $W_v$ ; #shards  $K$ ; mask ratio  $\rho$ ; thresholds  
978  $(\epsilon, \tau_E, \tau_{\text{prune}}, \tau_{\text{correct}}, \epsilon_{\text{cons}})$ ; margins  $(\gamma_1, \gamma_2, \gamma)$ ; KD weights  $(\lambda, \vartheta)$  for Eq.(4); routing-loss weight  $\lambda_a$ ;  
979 batch size  $b$ ; optional temperature  $T$  for soft KD.

980 1: Initialize side memories  $W'_{v,i} \leftarrow W_v$  and masks  $M_i \sim \text{Bernoulli}(\rho)$  for  $i = 1..K$ ; feedback pool  $\mathcal{E} \leftarrow \emptyset$ ;  
981 residual pool  $\mathcal{R} \leftarrow \emptyset$ .

982 2: **for** each incoming edit triple  $(x_e, y_e, x_{\text{loc}})$  **do**

983     3:      $i^* \leftarrow \text{ASSIGNSHARD}(x_e)$  ▷ Shard assignment by activation score

984     4:      $\mathcal{B} \leftarrow \text{FORMBATCHES}(\{x_e\} \cup \mathcal{R}, b)$  ▷ Distribution-aware batching

985     5:     **for** each batch  $B = \{x^{(0)}, \dots, x^{(b-1)}\} \in \mathcal{B}$  **do**

986         6:          $i \leftarrow \text{ASSIGNSHARD}(x^{(0)})$  ▷ Target shard for this batch

987         7:          $L_{\text{edit}} \leftarrow \text{AUTOREGCE}(B)$  ▷ Autoregressive cross-entropy

988         8:          $L_{\text{KD}} \leftarrow \text{INTRABATCHKD}(B, \lambda, \vartheta, T)$  ▷ Eq.(4); optional soft KD

989         9:          $L_{\text{act}} \leftarrow \text{ROUTINGMARGIN}(B, \gamma_1, \gamma_2, \gamma)$  ▷ Eq.(7)

990         10:         $L_{\text{batch}} \leftarrow L_{\text{edit}} + \lambda_a L_{\text{act}} + L_{\text{KD}}$

991         11:         $\text{MASKEDUPDATE}(W'_{v,i}, M_i, L_{\text{batch}})$  ▷  $W'_{v,i} \leftarrow W'_{v,i} - \eta(M_i \odot \nabla L)$

992         12:         $\text{FILTERANDRECLUSTER}(B, \epsilon_{\text{cons}}, \mathcal{R})$  ▷ Move high- $L_{\text{KD}}$  samples to residual pool

993         13:        **end for**

994         14:         $(\hat{y}, c) \leftarrow \text{EVALUATE}(x_e, y_e)$  ▷  $c \in \{0, 1\}$  indicates success

995         15:        **if**  $c = 0$  **then**

996             16:         $\mathcal{E} \leftarrow \mathcal{E} \cup \{(x_e, y_e)\}$

997             17:        **end if**

998         18:        **if**  $|\mathcal{E}| > \tau_E$  **or**  $\max_i \text{ERRORRATE}(\mathcal{E}, i) > \tau_{\text{prune}}$  **then**

999             19:         $\text{RETRIGGER}(\mathcal{E})$  ▷ Prune worst shard, rebuild, and retrain

1000         20:        **end if**

1001         21:        **end for**

1002         22:         $\text{LOSSAWARETIESMERGE}(\{W'_{v,i}\}_{i=1}^K, W_v)$  ▷ Loss-aware weighted TIES merge

---

1003  
1004  
1005  
1006  
1007  
1008  
1009

**Algorithm 2** REPAIR Inference with Dual-Memory Routing

---

1010 1: **function** ROUTEANDPREDICT( $x$ )

1011 2:     compute  $a(x) \leftarrow \text{FFNACTIVATION}(x)$  ▷ Activation  $A(x)$  at the target FFN layer

1012 3:     **for**  $i = 1..K$  **do**

1013         4:          $\Delta_{\text{act}}^{(i)}(x) \leftarrow \|a(x) \cdot (W'_{v,i} - W_v)\|_2$

1014         5:         **end for**

1015         6:         **if**  $\max_i \Delta_{\text{act}}^{(i)}(x) \leq \epsilon$  **then**

1016             7:             **return**  $f_{\theta_0}(x; W_v)$  ▷ Route to main memory

1017         8:         **else**

1018             9:              $i^* \leftarrow \arg \max_i \Delta_{\text{act}}^{(i)}(x)$

1019             10:             **return**  $f_{\theta_0}(x; W'_{v,i^*})$  ▷ Route to side memory  $i^*$

1020         11:         **end if**

1021         12:         **end function**

---

1022  
1023  
1024  
1025

---

1026  
1027  
1028  
1029  
1030   **Algorithm 3** Subroutines

---

1031   1: **function** ASSIGNSHARD( $x$ )  
1032   2:     $a \leftarrow \text{FFNACTIVATION}(x)$ ;  $\Delta^{(i)} \leftarrow \|a \cdot (W'_{v,i} - W_v)\|_2$ ,  $i = 1..K$   
1033   3:    **return**  $\arg \max_i \Delta^{(i)}$  ▷ Use the most active shard during training  
1034   4: **end function**  
1035   5: **function** FORMBATCHES( $S, b$ ) ▷ Distribution-aware batching  
1036   6:     $o_i \leftarrow \text{Norm}(\text{ModelFeat}(x^{(i)}))$  for  $i = 0, \dots, b - 1$   
1037   7:    Greedy seeding: pick  $x^{(0)} = \arg \max_{x \in S} \frac{1}{|S|} \sum_{x'} \cos(o(x), o(x'))$   
1038   8:    Build  $B \leftarrow \{x^{(0)}\} \cup \text{Top-}(b-1)$  nearest by cosine; remove  $B$  from  $S$   
1039   9:    Repeat until  $S$  is empty; **return** list of batches  $B$   
1040 10: **end function**  
1041 11: **function** AUTOREGCE( $B$ ) ▷ Autoregressive edit loss  $L_{\text{edit}}$   
1042 12:     $L \leftarrow 0$   
1043 13:    **for**  $x \in B$  with target sequence  $y$  **do**  
1044 14:       $L \leftarrow L - \sum_{t=1}^{|y|} \log p_\theta(y_t | y_{<t}, x)$   
1045 15:    **end for**  
1046 16:    **return**  $L/|B|$   
1047 17: **end function**  
1048 18: **function** INTRABATCHKD( $B, \lambda, \vartheta, T$ ) ▷ Eq.(4); optional soft-KD  
1049 19:    Compute  $o_i \leftarrow \text{Norm}(\text{ModelFeat}(x^{(i)}))$  for  $i = 0, \dots, b - 1$   
1050 20:     $L_{\text{cos}} \leftarrow \frac{1}{b-1} \sum_{i=1}^{b-1} \left( 1 - \frac{o_i^\top o_0}{\|o_i\| \|o_0\|} \right)$   
1051 21:     $o_{\text{mean}} \leftarrow \frac{1}{b} \sum_{i=0}^{b-1} o_i$ ;  $L_{\text{var}} \leftarrow \frac{1}{b} \sum_{i=0}^{b-1} \|o_i - o_{\text{mean}}\|_2^2$   
1052 22:     $L \leftarrow \lambda L_{\text{cos}} + \vartheta L_{\text{var}}$  ▷ Optional: KL distillation for added stability  
1053 23:    **if**  $T > 0$  **then**  
1054 24:      Get logits  $z_i$ ;  $p_i = \text{softmax}(z_i/T)$ ;  $L \leftarrow L + \frac{1}{b-1} \sum_{i=1}^{b-1} \text{KL}(p_0 \| p_i )$   
1055 25:    **end if**  
1056 26:    **return**  $L$   
1057 27: **end function**  
1058 28: **function** ROUTINGMARGIN( $B, \gamma_1, \gamma_2, \gamma$ ) ▷ Eq.(7)  
1059 29:     $L \leftarrow 0$   
1060 30:    **for** each edit sample  $x_e \in B$  **do**  
1061 31:      sample unrelated  $x_i$ ; compute  $\Delta_e = \text{ACTDELTA}(x_e)$ ,  $\Delta_i = \text{ACTDELTA}(x_i)$   
1062 32:       $L \leftarrow L + \max(0, \Delta_i - \gamma_1) + \max(0, \gamma_2 - \Delta_e) + \max(0, \gamma - (\Delta_e - \Delta_i))$   
1063 33:    **end for**  
1064 34:    **return**  $L/|B|$   
1065 35: **end function**  
1066 36: **function** MASKEDUPDATE( $W'_{v,i}, M_i, L$ ) ▷ Masked gradient to reduce cross-shard interference  
1067 37:     $g \leftarrow \nabla_{W'_{v,i}} L$ ;  $g_m \leftarrow M_i \odot g$  ▷  $M_i \in \{0, 1\}^{\text{shape}(W_v)}$   
1068 38:     $W'_{v,i} \leftarrow \text{OptimizerStep}(W'_{v,i}, g_m)$  ▷ SGD/Adam, etc.  
1069 39: **end function**  
1070 40: **function** FILTERANDRECLUSTER( $B, \epsilon_{\text{cons}}, \mathcal{R}$ )  
1071 41:    **for**  $x \in B$  **do**  
1072 42:       $\ell_{\text{KD}}(x) \leftarrow \text{per-sample KD vs. } x^{(0)}$   
1073 43:      **if**  $\ell_{\text{KD}}(x) \geq \epsilon_{\text{cons}}$  **then**  
1074 44:        move  $x$  to  $\mathcal{R}$   
1075 45:      **end if**  
1076 46:    **end for**  
1077 47:    **return**  
1078 48: **end function**  
1079 49: **function** EVALUATE( $x_e, y_e$ )  
1080 50:     $\hat{y} \leftarrow \text{ROUTEANDPREDICT}(x_e)$ ;  $c \leftarrow \mathbf{1}[\hat{y} = y_e]$   
1081 51:    **return**  $(\hat{y}, c)$   
1082 52: **end function**

---

---

1080  
1081  
1082  
1083  
1084  
1085  
1086  
1087  
1088  
1089

**Algorithm 4** Subroutines2

---

```

1: function ERRORRATE( $\mathcal{E}, i$ ) ▷ Error rate for shard  $i$ 
2:    $\mathcal{E}_i \leftarrow \{x \in \mathcal{E} \mid \arg \max_j \Delta_{\text{act}}^{(j)}(x) = i\}$ 
3:    $r_i \leftarrow \frac{|\{x \in \mathcal{E}_i \mid \text{CORRECTNESS}(x) \leq \tau_{\text{correct}}\}|}{|\mathcal{E}_i|}$ 
4:   return  $r_i$ 
5: end function
6: procedure RETRIGGER( $\mathcal{E}$ ) ▷ Closed-loop pruning and retraining
7:    $j \leftarrow \arg \max_i \text{ERRORRATE}(\mathcal{E}, i)$  ▷ Identify worst-performing shard
8:   Remove or reinitialize shard  $j$ :  $W'_{v,j} \leftarrow W_v + \sigma_{\text{init}} \cdot \mathcal{N}(0, 1)$ ; resample  $M_j$ 
9:   Build  $\mathcal{E}_{\text{retrain}}$  from  $\mathcal{E}$ ; form batches; retrain shards via MASKEDUPDATE + INTRABATCHKD
10: end procedure
11: function LOSSAWARETIESMERGE( $\{W'_{v,i}\}, W_v$ ) ▷ Loss-aware weighted TIES merge
12:   For each shard  $i$ :  $\tau_i \leftarrow W'_{v,i} - W_v$ ; compute training loss  $L_i$  on its assigned data
13:    $w_i \leftarrow \frac{e^{-\alpha L_i}}{\sum_j e^{-\alpha L_j}}$ 
14:   for each parameter index  $p$  do
15:      $S \leftarrow \{(i, \tau_i[p], w_i)\}_{i=1}^K$ 
16:     if all  $\tau_i[p]$  share the same sign then
17:        $\delta[p] \leftarrow \sum_i w_i \tau_i[p]$  ▷ Consistent signs: weighted sum
18:     else
19:        $i^* \leftarrow \arg \max_i \{w_i \mid \tau_i[p]\}$ ;  $\delta[p] \leftarrow \tau_{i^*}[p]$  ▷ Conflict: keep most trustworthy shard
20:     end if
21:   end for
22:    $W_v \leftarrow W_v + \delta$ ; return  $W_v$ 
23: end function
24: function FFNACTIVATION( $x$ )
25:   return activation  $A(x)$  at the target FFN layer
26: end function
27: function ACTDELTA( $x$ )
28:   return  $\max_i \|A(x) \cdot (W'_{v,i} - W_v)\|_2$ 
29: end function
30: function MODELFEAT( $x$ )
31:   return feature used for similarity (e.g.,  $A(x)$  or last-token state)
32: end function
33: function NORM( $v$ )
34:   return  $v / \|v\|_2$ 
35: end function
36: function CORRECTNESS( $x$ )
37:   return predicted correctness score for  $x$ 
38: end function

```

---

1126  
1127  
1128  
1129  
1130  
1131  
1132  
1133